ACT II · TRACK 1

General-Purpose AI Agent.

WAYNE ◆ JERMAINE

batch inference agent

● Docker · linux/amd64 ● Python ● Fireworks AI

THE OBJECTIVE

Enough accuracy. Minimum tokens.

Not a chatbot — a **batch agent**. Read every prompt from `/input/tasks.json`, solve it, write `/output/results.json`, exit 0. No frontend, no user.

STAGE 1

$\geq$ 80%

accuracy gate · 16 of 19 tasks

below the gate → off the leaderboard

STAGE 2

Fewest tokens

token efficiency · passers only

ranked ascending → fewest wins

8 CAPABILITY AREAS

factual | math | sentiment | summarisation | NER | debugging | logic | code gen

THE AGENT

Answer locally. Escalate rarely.

Only tokens through `FIREWORKS_BASE_URL` are scored.
Every verified local answer is free.

● Docker · linux/amd64 ● Python ● Fireworks AI

IMPLEMENTATION · ITERATION

We guessed. The benchmark disagreed.

Every number below is measured, not assumed.

01	**Hidden reasoning ate the token cap** thinking filled the cap, answers came back empty	**–27%** with reasoning_effort = low
02	**Routing between models wasn't the lever** optimal per-category routing vs one frontier model	**12%** total token saving
03	**Our hypothesis was backwards** we predicted math, logic and code would need the API	**100%** local score on those three

● Docker · linux/amd64 ● Python ● Fireworks AI

THE RESULT

90% accuracy. Zero tokens.

The bundled 3B model clears all eight categories alone — 36 of 40 correct. Fireworks is the fallback, not the default.

measured on our own 40-task harness
the live gate is an LLM judge over 19 unseen tasks

5 TASKS PER CATEGORY

TRACK 1 · SUBMISSION

Thank you.

A 3B model, a verifier, and an API we barely call.

WAYNE ◆ JERMAINE

github.com/WayneCh0y/amd-hack